BENJAMIN TOUSIFAR

San Diego, CA • (619) 643-6899 • bn.tousifar86@gmail.com • github.com/BeTous8 • linkedin.com/in/benjamin-tousifar

SUMMARY

Data Engineer with 2+ years of production experience operating ETL pipelines for Fortune 500 clients, processing millions of daily transactions across SAP HANA, Oracle, and Azure. Certified in Azure Data Engineering (DP-203) and Snowflake (SnowPro Core). Strong SQL and Python with hands-on data quality, reconciliation, and Change Data Capture experience in enterprise environments. B.S. Electrical and Computer Engineering, UC San Diego.

TECHNICAL SKILLS

Data Engineering: ETL/ELT Pipeline Development, Informatica PowerCenter, Azure Data Factory, Change Data Capture (CDC), Incremental Loading, Data Reconciliation, Data Quality Validation, Multi-Source Integration
Databases & SQL: PostgreSQL, Snowflake, Oracle, SQL Server, MySQL, Hive, SAP HANA, Schema Design, Query Optimization, Window Functions, CTEs, Stored Procedures
Cloud & Warehousing: Azure Data Factory, Azure Synapse Analytics, Azure Blob Storage, Snowflake, Supabase
Programming: Python (Pandas, SQLAlchemy, Pydantic, FastAPI), SQL, JavaScript, TypeScript, Node.js
Engineering Practices: Git, Bitbucket, Azure DevOps, CI/CD, Unit Testing, Dev/QA/Prod Deployment, Agile/Scrum, Jira, Confluence, ServiceNow
Certifications: Azure Data Engineer Associate (DP-203), Snowflake SnowPro Core, Azure Fundamentals (AZ-900)

DATA ENGINEERING EXPERIENCE

Data Engineer — LTIMindtree — Jan 2022 - Mar 2024

Honeywell Account (Dec 2022 - Jan 2024)

- Operated and maintained production ETL pipelines in Informatica PowerCenter integrating SAP HANA, Oracle, and third-party sources into a cloud data warehouse, processing millions of daily transactions across 7-8 Strategic Business Units
- Investigated and resolved data quality incidents by writing SQL and Python reconciliation checks against both source and target systems to isolate divergent records and identify root cause; tracked and closed incidents within SLA through ServiceNow
- Operated Change Data Capture (CDC) pipelines loading only modified records via insert/update/delete operations, and diagnosed failure modes including late-arriving and out-of-order changes that silently corrupt target data

Johnson Controls Account (Mar 2022 - Oct 2022)

- Maintained and modified ETL pipelines in Azure Data Factory integrating ERP data into Hive and Azure Synapse Analytics, promoting SQL transformation changes across Dev/QA/Prod via Azure DevOps and Bitbucket and validating output against known results before each promotion

DATA-INTENSIVE PROJECTS

Memora — Founder & Engineer (memoraapp.netlify.app) — 2024 - Present

- Modeled a normalized PostgreSQL schema (Supabase) where correctness is enforced by the database rather than the client: foreign keys with deliberate ON DELETE semantics (RESTRICT on any row carrying financial history), CHECK constraints on monetary and enum columns, triggers that freeze item fields mid-fulfillment, and a partial unique index permitting only one active fulfillment per item
- Built an ingestion pipeline for heterogeneous third-party product data: canonicalizes URLs to stable product keys (ASIN / TCIN / SKU), prefers structured APIs over scraping, then falls back through JSON-LD → embedded JSON → Open Graph → per-store selectors, coalescing per field by source trust and emitting completeness metadata so partial records degrade instead of failing
- Made payment-event processing idempotent under at-least-once webhook delivery: unique constraint on the provider session id, ON CONFLICT DO NOTHING inserts, atomic in-place increments, and SERIALIZABLE isolation with SELECT ... FOR UPDATE on redemption writes behind a per-request idempotency key

Lender Matching Platform — Engineering Assessment (GitHub) — 2024

- Built a metadata-driven rules engine over PostgreSQL where underwriting policies are stored as data rather than code, so onboarding a new lender requires only new rows and no deployment
- Modeled entities as SQLAlchemy ORM classes with foreign key relationships, enforced schema contracts on all inputs and outputs with Pydantic, and wrote 25 passing unit tests covering business logic and edge cases
- Diagnosed and resolved an N+1 query pattern using selectinload eager loading, collapsing per-record lookups into a single batched fetch
- **SQL Analytics Portfolio** — PostgreSQL schema with 5 tables and 20+ analytical queries: window functions (DENSE_RANK, NTILE, LAG/LEAD), recursive CTEs, GROUPING SETS, and business intelligence reporting patterns (GitHub)

EDUCATION

University of California, San Diego — B.S. Electrical and Computer Engineering — June 2021

Relevant Coursework: Data Structures, Algorithms, Database Systems, Object-Oriented Programming, Software Engineering